

Week 7: Monte Carlo (MC)

MATH-517 Statistical Computation and Visualization

Linda Mhalla

2025-10-31

Introduction

MC $\equiv$ repeated random sampling to mimic outcome of random process and produce numerical results such as

- generating draws from complicated distributions and/or domains
- integration
 - calculation of moments or confidence intervals
 - high-dimensional densities in Bayesian settings
- optimization
 - mode evaluation

Basic idea: If we can sample from a process or mimic its outcomes, we can learn a lot about it by doing statistics on the simulated samples (as opposed to analyzing the process itself)

MC methods $\equiv$ simulation-based statistical techniques/inference

Introduction

Gambling experiments have random outcomes – hence “Monte Carlo”



Method initially developed by Stanislaw Ulam and John von Neumann for the Manhattan Project (to estimate integrals)

Example

$(X_1, Y_1)^\top, \dots, (X_N, Y_N)^\top$ a sample from the standardized bivariate Gaussian distribution

$$\mathcal{N} \left(\begin{pmatrix} 0 \\ 0 \end{pmatrix}, \begin{pmatrix} 1 & \rho \\ \rho & 1 \end{pmatrix} \right)$$

- We want to test $H_0 : \rho = \rho_0$ against $H_1 : \rho \neq \rho_0$
- Statistic: $\hat{\rho}_N = N^{-1} \sum_{n=1}^N X_n Y_n$

Since the data generation process is fully determined under H_0 , we can simulate data to approximate the sampling distribution and thus also the p -value

To test a hypothesis, we only need to simulate data under H_0

But how to draw samples from a specific distribution?

Example: Estimation of π

$$\Pr(\text{point inside the circle}) = \frac{\text{area of the circle}}{\text{area of the square}} = \frac{\pi r^2}{2r \times 2r} = \frac{\pi}{4}$$

```
set.seed(123)
piVal <- function(nPoints = 300){
  x <- runif(nPoints,-1,1)
  y <- runif(nPoints,-1,1)
  result <- ifelse(x^2+y^2 <= 1, TRUE, FALSE)
  4*sum(result)/nPoints
}

N <- 1000
pi_est <- replicate(N, piVal())
mean(pi_est) #to get uncertainty
```

```
[1] 3.141893
```

Section 1

Random Number Generation (RNG)

RNG

True randomness is hard to come by. Historically:

- dice, cards, coins
- physical processes
- census data, tables, etc.

Practical reasons not to use "truly" random numbers:
debugging and reproducibility

John von Neumann: pseudo-RNG

- approximates the desired dist. for $N \rightarrow \infty$
- cannot be predicted (without the seed or generator)
- pass a set of independence tests
- repeatability ($\Rightarrow$ reproducibility)
- long cycle (before it starts repeating) and fast sampling



Uniformity and independence tests needed to assess quality of pseudo-RNG

Cornerstone: Generating from $\mathcal{U}[0, 1]$

Assume now we can generate numbers from the $\mathcal{U}[0, 1]$ distribution

- e.g., the *linear congruential method (LCG)*

$$X_n = (aX_{n-1} + c) \bmod m, \quad n = 1, 2, \dots,$$

where a, c, m , and X_0 are cleverly chosen to fulfill the pseudo-RNG requirements, i.e., maximize period, speed, and “randomness”

- X_0 is the seed
- produces integers between 0 and $m - 1$
- $U_n = X_n/m \sim U(0, 1)$

Bad example: $m = 2^{31}$, $a = 2^{16} + 3$, and $c = 0 \Rightarrow$ IBM's RANDU

Example: $a = 13$, $c = 0$, and $m = 64$

- $X_0 = 4$: 4, 52, 36, 20, 4
- $X_0 = 2$: 2, 26, 18, 42, 34, 58, 50, 10, 2

Choice of parameters

Case	Parameters	Max period	Conditions
Mixed LCG	$m = 2^b, c \neq 0$	$P = 2^b$	c relatively prime to m ; $a = 1 + 4k, k \in \mathbb{Z}$
Multiplicative LCG	$m = 2^b, c = 0$	$P = 2^{b-2}$	Seed X_0 odd; $a = 3 + 8k$ or $a = 5 + 8k, k = 0, 1, \dots$
Multiplicative (prime modulus)	m prime, $c = 0$	$P = m - 1$	a must be a <i>primitive root</i> mod m (smallest k with $a^k \equiv 1 \pmod{m}$ is $k = m - 1$)

Now, better and much more complicated algorithms are available

- every piece of software has its favorite pseudo-RNG
- outside the scope of the course (see, e.g., shift-register generators or [Wichmann–Hill generator](#))

Question: How do we generate from other distributions?


Transforms

Lemma. (*Inverse Transform.*)

Let $U \sim \mathcal{U}(0, 1)$ and F be a distribution function and F^{-1} the quantile (or generalized inverse) function. Then $X = F^{-1}(U) \sim F$.

Proof: Simply

$$P(X \leq x) = P\{F(X) \leq F(x)\} = P\{U \leq F(x)\} = F(x).$$



The inverse transform method is general, but requires a known and easily invertible CDF, which many important distributions don't have in closed form

- e.g., $\mathcal{N}(0, 1)$

Often, simpler relationships can be used: [diagram](#)

- still, there is no arrow there between $\mathcal{U}(0, 1)$ and $\mathcal{N}(0, 1)$, generating $\mathcal{N}(0, 1)$ is actually a bit tricky...

Transforms

Lemma. (*Box-Muler transform.*)

Let $U_1, U_2 \sim \mathcal{U}(0, 1)$ be independent. Then

$$Z_1 = \sqrt{-2 \log(U_1)} \cos(2\pi U_2) \quad \& \quad Z_2 = \sqrt{-2 \log(U_1)} \sin(2\pi U_2)$$

are two independent standard Gaussian random variables

Explanation:

- points from the $\mathcal{N}(0, I_2)$, when written in polar coordinates, have an angle $\theta \sim U[0, 2\pi)$ which is independent of the radius R
- the squared radius $R^2 = Z_1^2 + Z_2^2 \sim \chi_2^2$, i.e. $\text{Exp}(1/2)$

Again, software uses its favorite relationships

- e.g., R has tabulated F and F^{-1} for $\mathcal{N}(0, 1)$ to a high precision and actually uses the inverse transform, because evaluating trigonometric functions is rather expensive (slow). `?rnorm` $\Rightarrow$ `rnorm`, `pnorm`, `qnorm`, `dnorm` help
- NumPy uses the [Ziggurat algorithm](#) (faster than the others)

Rejection Sampling

Setup: we know how to simulate from a *proposal* g , we want to simulate from a *target* f


- let $\text{supp}(f) \subset \text{supp}(g)$, i.e., $f(x) > 0 \Rightarrow g(x) > 0$
- let there be $c > 1$ such that $\forall x : f(x) \leq c g(x)$, i.e., $\sup_x \frac{f(x)}{g(x)} = c < \infty$

Algorithm: (to obtain a single draw X from f)

- 1 Draw a proposal Y from g
- 2 Draw $U \sim \mathcal{U}(0, 1)$
- 3 If $U \leq \frac{1}{c} \frac{f(Y)}{g(Y)}$, accept $X = Y$ and stop, otherwise go back to 1

Example:

- $\mathcal{U}(0, 1)$ proposal
- $\mathcal{B}(2, 5)$ target
- $c \approx 2.5$



Rejection Sampling

Does the algorithm really sample from f ?

$$\begin{aligned} P(X \leq x) &= P\left\{Y \leq x \mid U \leq \underbrace{\frac{1}{c} \frac{f(Y)}{g(Y)}}_{=:t(Y)}\right\} = \frac{P\{Y \leq x \wedge U \leq t(Y)\}}{P\{U \leq t(Y)\}} \\ &= \frac{\int_{-\infty}^x \int_0^{t(y)} du g(y) dy}{\int_{-\infty}^{+\infty} \int_0^{t(y)} du g(y) dy} = \frac{\int_{-\infty}^x t(y) g(y) dy}{\int_{-\infty}^{+\infty} t(y) g(y) dy} = \frac{\int_{-\infty}^x \frac{1}{c} f(y) dy}{\int_{-\infty}^{+\infty} \frac{1}{c} f(y) dy} \\ &= \frac{\frac{1}{c} F(x)}{\frac{1}{c}} = F(x) \end{aligned}$$


The rejection sampling algorithm above is again quite general, but it needs

- a good proposal g
 - as close as possible to the target density (small values of c), leading to
 - high acceptance probability $P\{U \leq t(Y)\} = 1/c$
- fast evaluation of f and g

Example: $\mathcal{N}(0, 1)$ again

Goal: Simulate data from the standard Gaussian target using the double exponential (Laplace) proposal, i.e.,

$$f(x) = \frac{1}{\sqrt{2\pi}} e^{-\frac{x^2}{2}} \quad \& \quad g(x) = \frac{\alpha}{2} e^{-\alpha|x|}, \text{ where } \alpha > 0, \quad x \in \mathbb{R}$$



Another way of obtaining
 $\mathcal{N}(0, 1)$ from $\mathcal{U}(0, 1)$:

$$\mathcal{U}(0, 1) \longrightarrow \text{Exp}(1)$$

$$\text{Exp}(1) \longrightarrow \text{DbExp}(1)$$

$$\text{DbExp}(1) \longrightarrow \mathcal{N}(0, 1)$$

Note: $\alpha = 1$ minimizes the value of $c = \sup f(x)/g(x)$ and hence maximizes the acceptance probability

Section 2

Numerical Integration




Deterministic Quadrature Approaches

Goal: approximate $J = \int_a^b f(x)dx$ via composite quadrature rules

Quadrature method: evaluate the function on a grid

$$S_K = \{(b-a)/K\} \sum_{k=1}^K f(t_k)$$

- if f is nice (smooth), $S_K \rightarrow J$ for $K \rightarrow \infty$
- corresponds to integrating a local constant interpolation of f
- local linear interpolation (*trapezoidal rule*) or local quadratic (*Simpson's rule*) are also well known



Naive Monte Carlo

We consider the more general integral

$$J = \int_{\mathcal{X}} m(x)f(x)dx = \mathbb{E}_f\{m(X)\} \quad \text{for } X \sim f$$

$\Rightarrow$ generate $X_1, \dots, X_N \stackrel{i.i.d.}{\sim} f$ and approximate J by

$$\widehat{J}_N = N^{-1} \sum_{n=1}^N m(X_n)$$

- Provided $\int |m(x)|f(x)dx < \infty$, unbiased and consistent due to SLLN
- monitoring convergence via CLT-based confidence intervals:

$$\sqrt{N} \frac{\widehat{J}_N - J}{\sigma(m)} \sim \mathcal{N}(0, 1), \quad \text{where } \sigma^2(m) = \int m^2(x)f(x)dx - J^2$$

(Naive Monte Carlo estimate v_N^2 of $\sigma^2(m)$ is used in practice)

$\Rightarrow$ MC integration methods give a probabilistic error bound versus a deterministic one with numerical quadrature rules

Naive Monte Carlo

$$\lim_{N \rightarrow \infty} \Pr \left(-3 \frac{\sigma(m)}{\sqrt{N}} \leq \frac{1}{N} \sum_{n=1}^N m(X_n) - J \leq 3 \frac{\sigma(m)}{\sqrt{N}} \right) \approx 99.73\%$$

$\Rightarrow$ to attain an error of $\epsilon > 0$ with 99.73%, we need

$$N = \frac{9}{\epsilon^2} \sigma^2(m) = O(1/\epsilon^2)$$

- This is independent of the dimension (depends on the smoothness of m though)!
- Deterministic numerical integration methods suffer the curse of dimensionality: $N = O(1/\epsilon^d)$, roughly
- Monte Carlo estimate for the integral converges slowly to the true value (slower than quadrature methods for smooth functions in low dimensions)
- Beware of **rare events**: if f has heavy tails, $\widehat{J}_N$ can be a bad estimate and we need huge N to get small $v_N \Rightarrow$ techniques for variance reduction

Importance Sampling

We often can not simulate directly from f and we require sophisticated approaches. Rewrite

$$J := \int_{\mathcal{X}} m(x)f(x)dx = \int_{\mathcal{X}} m(x)\frac{f(x)}{g(x)}g(x)dx = \int_{\mathcal{X}} m(x)w(x)g(x)dx$$

with g a density whose support contains that of f and $w(x) \geq 0$ the **importance weighting function**

Thus, by sampling $X_1, \dots, X_N \stackrel{\text{i.i.d.}}{\sim} g$, we can approximate J by

$$\hat{J}_N := N^{-1} \sum_{n=1}^N m(X_n)w(X_n)$$

Idea:

- Use a simpler proposal distribution g from which we can generate
- Candidates generated from g fall within the domain of f
- Reweight the observations generated from it when taking the mean

Importance Sampling: Intuitive Explanation

Key: integrating f amounts to integrating f/g under sampling from g

$$J = \int_{\mathcal{X}} f(x) dx = \int_{\mathcal{X}} m(x) f(x) g(x) dx = \mathbb{E}_g\{m(X)w(X)\}$$

- when f is flat (all regions are equally important), use either the naive MC (with uniform sample) or deterministic approaches that need only small samples
- when f is not flat, using a “good” g allows us to encode which regions are important $\Rightarrow$ “importance sampling” (vs rejection sampling)

Of course, it is not always easy to find a “good” g which

- has a similar shape than f and
- from which we can easily sample

As we will see, when $\mathcal{X} = \mathbb{R}$, it is important to match the decay of the tails between the target and reference measures

Importance Sampling: Properties

- unbiased and the variance is given by

$$\begin{aligned}\text{var}(\hat{J}_N) &= \frac{1}{N} \left\{ \int_{\mathcal{X}} m^2(x) \frac{f(x)}{g(x)} f(x) dx - J^2 \right\} \\ &= N^{-1} \int_{\mathcal{X}} \left\{ m(x) \frac{f(x)}{g(x)} - J \right\} m(x) f(x) dx,\end{aligned}$$

which is small if $g(x) \approx m(x)f(x)J^{-1}$ or $g(x) \propto m(x)f(x)$

$\Rightarrow$ good choices of g can yield huge improvements in efficiency, e.g.,

- g with similar shape than mf
- maximum principle ($g(x)$ and $m(x)f(x)$ take their maximum at the same value)
- g from the same distribution family as f

For instance, reduction in variance is

$$\text{var}(\hat{\hat{J}}_N) - \text{var}(\hat{J}_N) = \frac{1}{N} \int_{\mathcal{X}} m^2(x) \left\{ 1 - \frac{f(x)}{g(x)} \right\} f(x) dx$$


Importance Sampling: Properties

- Most difficult aspect to importance sampling is in choosing a good sampling density g
- Need to be very careful as it is possible to choose g according to some heuristics, but that results in a variance increase
- It is possible to have an importance sampling estimator with infinite variance
 - e.g., if g puts too little weight relative to f on the tails of the distribution

Examples

Task: Approximately calculate $P(2 < X < 6)$ for the target distribution $X \sim f$ using a reference g


Gaussian target, Exponential reference - densities (left), their ratio (middle), importance sampling error (right)



Examples

Approximately calculate $P(2 < X < 6)$ for the target distribution $X \sim f$ using a reference g


Cauchy target, Exponential reference - densities (left), their ratio (middle), importance sampling error (right)



Examples

Approximately calculate $P(2 < X < 6)$ for the target distribution $X \sim f$ using a reference g

Cauchy target, Gaussian reference - densities (left), their ratio (middle), importance sampling error (right)



- the tails of Cauchy and Gaussian distributions are too different $\Rightarrow$ importance sampling performs poorly
- if we can simulate from Gaussian, we can simulate directly from Cauchy: $Z_1, Z_2 \sim \mathcal{N}(0, 1)$ independent $\Rightarrow Z_1/Z_2 \sim \text{Cauchy}(0, 1)$

Variance Reduction

Accuracy of MC integration is assessed by the estimator's efficiency/variance (assuming efforts of simulation are similar)

There are ways to tweak the sampling scheme in order to reduce the variance

- importance sampling (we have seen above)
- antithetic variables (to follow)
- stratified sampling (to follow)
- quasi-random sampling and control variates (see the [supplementary notes](#))
- many other techniques: latin hypercube sampling, ratio estimator, etc

Remark: When comparing several different estimators via simulations, the same simulated datasets should be used for all the estimators

Variance Reduction: Antithetic Variables

Idea: Introduce negative correlation between pairs of replications and rely on

$$\text{var}(f_1 + f_2) = \text{var}(f_1) + \text{var}(f_2) + 2\text{cov}(f_1, f_2)$$

Given two i.i.d. samples $X_1, \dots, X_N \sim f$ and $Y_1, \dots, Y_N \sim f$, consider the estimator

$$\tilde{J}_N = \frac{1}{2N} \sum_{n=1}^N \{m(X_n) + m(Y_n)\} = \frac{1}{2}(\hat{J}_N^X + \hat{J}_N^Y)$$

Then,

$$\text{var}(\tilde{J}_N) = \frac{1}{2} \text{var}(\hat{J}_N) \{1 + \text{corr}(\hat{J}_N^X, \hat{J}_N^Y)\}$$

$\Rightarrow \tilde{J}_N$ is more efficient than the naive MC (with sample of size $2N$) if $m(X_n)$ and $m(Y_n)$ are **negatively correlated**

Basic result: if $g(u)$ is monotonic on $0 < u < 1$, then

$$\text{corr}\{g(U), g(1 - U)\} < 0$$

Thus $F^{-1}(U)$ and $F^{-1}(1 - U)$ are negatively correlated with distribution F

Variance Reduction: Stratified Sampling

- Break sampling space into strata $\Delta_1, \dots, \Delta_k$ and sample appropriate number of observations in each
- Compute the naive MC estimator in each stratum and sum over all strata
- Beneficial when $m(X)$ (or $f(x)$) varies a lot across the space

We want to estimate $J = \mathbb{E}\{m(X)\}$. Let W be another r.v. s.t.

- $p_i = \Pr(W \in \Delta_i)$, with $\Delta_i \subset \mathbb{R}$, is easily computed
- we know how to generate $X \mid W \in \Delta_i$

Assuming $\Delta_1, \dots, \Delta_k$ is a partition of $\mathbb{R}$ and denoting $J_j = \mathbb{E}\{m(X) \mid W \in \Delta_j\}$ and $\sigma_j^2 = \text{var}\{m(X) \mid W \in \Delta_j\}$, we have

$$J = \mathbb{E}[\mathbb{E}\{m(X) \mid W\}] = \sum_{j=1}^k p_j J_j, \quad \text{and} \quad \hat{J}_{str,N} = \sum_{j=1}^k p_j \hat{J}_{j,N_j}$$

Variance Reduction: Stratified Sampling

- Method relies on conditional variance:

$$\text{var}\{m(X)\} = \mathbb{E}[\text{var}\{m(X)|W\}] + \text{var}[\mathbb{E}\{m(X) | W\}]$$

Thus,

$$\text{var}\{m(X)\} \geq \mathbb{E}[\text{var}\{m(X)|W\}] = \sum_{j=1}^k p_j \sigma_j^2$$

For instance, assuming proportional allocation, i.e., $N_j = Np_j$, then

$$\text{var}(\hat{J}_{str,N}) = \sum_{j=1}^k p_j^2 \frac{\sigma_j^2}{N_j} = \frac{\sum_{j=1}^k p_j \sigma_j^2}{N} \leq \text{var}(\hat{J}_N) = \frac{\text{var}\{m(X)\}}{N}$$

- Variance reduction substantial if W accounts for a large fraction of the variance of $m(X)$
- Variance reduction depends on the allocation in each stratum. Thus, we can minimize $\text{var}(\hat{J}_{str,N})$ subject to $\sum_{j=1}^m N_j = N \Rightarrow$ more samples go to strata with high variability $N_j \propto p_j \sigma_j$

Donald Knuth (1997, 3rd ed.) *The Art of Computer Programming*, vol. 2

Robert & Casella (2010) *Introducing Monte Carlo methods with R*

Assignment

Go to [Assignment 6](#) for details.